

# Discussion

CS 5/7320  
Artificial  
Intelligence

Intelligent Agents  
AIMA Chapter 2

---

Slides by Michael Hahsler  
with figures from the AIMA textbook.



This work is licensed under a [Creative Commons Attribution-ShareAlike 4.0 International License](https://creativecommons.org/licenses/by-sa/4.0/).





Image: "Robot at the British Library Science Fiction Exhibition"  
by BadgerGravling

A white humanoid robot, ASIMO, stands in a dimly lit living room. The robot is wearing a white hood and has "ASIMO" printed on its chest. In the background, there is a television displaying the word "ASIMO", a potted plant on a wooden table, and a window with curtains. The scene is set in a home environment.

# Module Review

# Designing a Rational Agent




Remember the definition of a rational agent:

*“For each possible percept sequence, a rational agent should select an **action** that **maximizes its expected performance measure**, given the evidence provided by the **percept sequence** and the **agent’s built-in knowledge**.”*

**Agent Function**

- Represents the “brain”
- Assess performance measure
- Remember percept sequence
- Built-in knowledge



**Important:** Everything outside the agent function represents the environment. This includes the physical robot, its sensors and its actuators, and event loop!

# Rational Agents

**Rule:** Pick the action that maximize the expected utility


$$a = \operatorname{argmax}_{a \in A} E(U \mid a)$$

This means:

- **Rationality is an ideal** – it implies that no one can build a better agent
- **Rationality  $\neq$  Omniscience** – rational agents can make mistakes if percepts and knowledge do not suffice to make a good decision
- **Rationality  $\neq$  Perfection** – rational agents maximize **expected** outcomes not actual outcomes
- **It is rational to explore and learn** – i.e., use **percepts** to supplement prior knowledge and become autonomous
- **Rationality is often bounded** by available memory, computational power, available sensors, etc.

# Simple Reflex Agent

- Uses only built-in knowledge in the form of **rules** that select action only **based on the current percept**. This is typically very fast!
- The **agent does not know about the performance measure**! But well-designed rules can lead to good performance.
- The agent needs no memory and ignores all past percepts.




The interaction is a sequence:  $p_0, a_0, p_1, a_1, p_2, a_2, \dots, p_t, a_t, \dots$

**Example:** A simple vacuum cleaner that uses rules based on its current sensor input.

# Model-based Reflex Agent

- Maintains a **state variable** to keep track of aspects of the environment that cannot be currently observed. I.e., it has memory.
- It knows how the environment evolves over time given its last action. It updates the state using a **transition function** and the new percept.
- There is now more information for the **rules** to make better decisions.

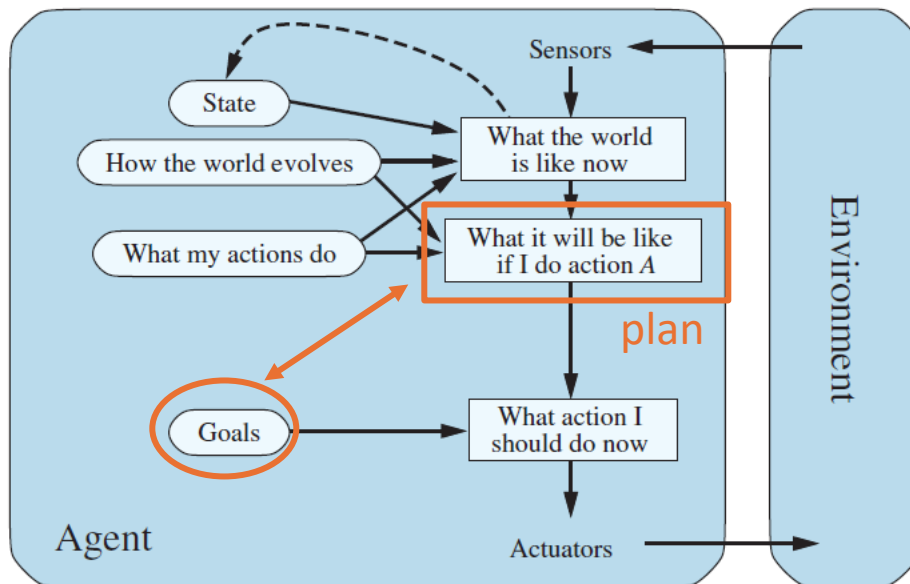


The interaction is a sequence:  $p_0, s_0, a_0, p_1, s_1, a_1, p_2, s_2, a_2, p_3, \dots, p_t, s_t, a_t, \dots$

**Example:** A vacuum cleaner that remembers where it has already cleaned.

# Goal-based Agent

- The agent has the task of reaching a defined **goal state** and is then finished.
- The agent needs to move towards the goal. As special type is a **planning agent** that uses **search algorithms** to plan a sequence of actions that leads to the goal.
- Performance measure: the **cost to reach the goal**.



$$a = \operatorname{argmin}_{a_0 \in A} \left[ \underbrace{\sum_{t=0}^T c_t}_{\text{Sum of the cost of a planned sequence of actions that leads to a goal state}} \mid s_T \in S^{\text{goal}} \right]$$

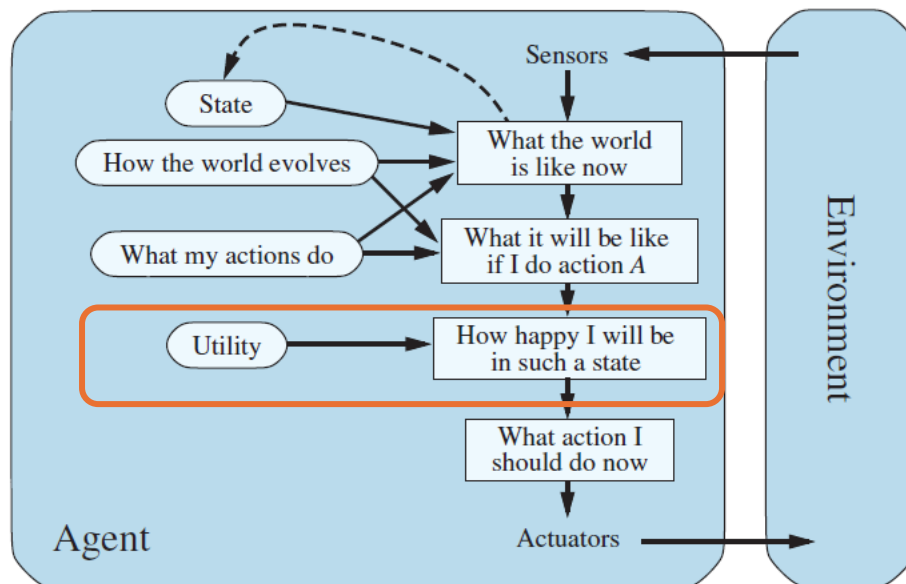
The interaction is a sequence:  $p_0, s_0, a_0, p_1, s_1, a_1, p_2, s_2, a_2, \dots, s^{\text{goal}}$

cost

**Example:** Solving a puzzle. What action gets me closer to the solution?

# Utility-based Agent

- The agent uses a utility function to evaluate the **desirability of each possible states**. This is typically expressed as the reward of being in a state  $R(s)$ .
- Choose actions to stay in desirable states.
- Performance measure: The discounted sum of **expected utility over time**.



$$a = \operatorname{argmax}_{a_0 \in A} E \left[ \underbrace{\sum_{t=0}^{\infty} \gamma^t r_t}_{\text{Implements rational behavior: Utility is the expected future discounted reward}} \mid a_0 \right]$$

Implements rational behavior: Utility is the expected future discounted reward

**Techniques:** Markov decision processes, reinforcement learning

The interaction is a sequence:  $p_0, s_0, a_0, p_1, s_1, a_1, p_2, s_2, a_2, \dots$

└────────────────────────────────┘  
reward

**Example:** An autonomous Mars rover prefers states where its battery is not critically low.



# Some Environment Types Revisited

**Fully observable:** The agent's sensors always show the whole **state**.

**vs.**

**Partially observable:** The agent only perceives part of the **state** and needs to remember or infer the rest.

**Deterministic:**

- a) **Percepts** are 100% reliable.
- b) Changes in the environment are completely determined by the current **state** of the environment and the agent's **action**.

**vs.**

**Stochastic:**


- a) **Percepts** are unreliable (noise distribution, sensor failure probability, etc.). This is called a stochastic sensor model.
- b) The **transition function** is stochastic leading to transition probabilities and a Markov process.

**Known:** The agent knows the **transition function**.

**vs.**

**Unknown:** The needs to **learn the transition function** by trying actions.

We will spend the whole course on discussing algorithms that can deal with environments that have different combinations of these three properties.

The background is a dark, abstract composition featuring horizontal motion blur streaks in shades of blue and orange. In the lower right, there are curved, concentric lines that suggest a road or a tunnel, creating a sense of depth and movement.

# Case Study: Self-Driving Cars



# Self-driving Cars

## SAE Automation Levels

- Level 1 - Driver Assistance (“hands on”)
- Level 2 - Partial Automation (“hands off”)
- Level 3 - Conditional Automation
- Level 4 - High Automation
- Level 5 - Full Automation (“steering wheel optional”)

## Components

- Sensing
- Maps
- Path planning
- Controlling the vehicle

## Why is this so hard?





# A Self-Driving Car as a Rational Agents

**Rule:** Pick the action that maximize the expected utility

$$a = \operatorname{argmax}_{a \in A} E(U \mid a)$$

Answer the following questions:

- If we have two cars and one provides more (expected) utility. Which car is rational?
- Can a rational self-driving car be involved in an accident?
- How would a self-driving car explore and learn?
- What does bounded rationality mean for a self-driving car?

# PEAS Description of the Environment of a Self-Driving Car



Complete the PEAS description.

Performance measure	Environment	Actuators	Sensors

# Percepts and States: Self-Driving Car



Describe percepts and states.

**Percepts**

**States**

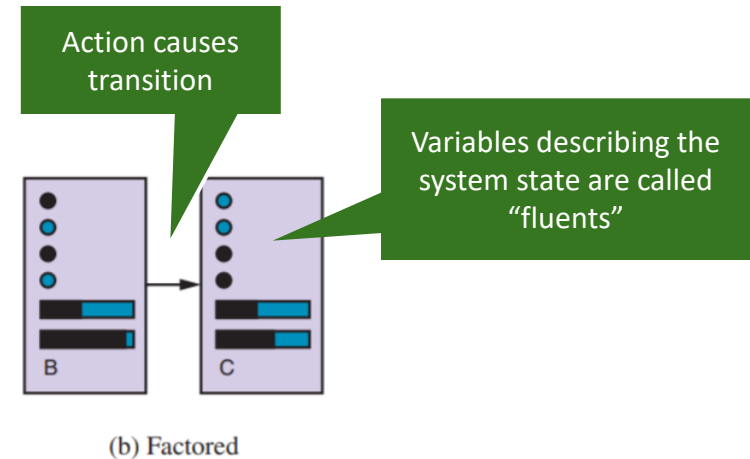


# State Representation: Self-Driving Car

States help to keep track of the environment and the agent in the environment.

Design a structured representation for the state of a self-driving car.

- a) What fluents should it contain?
- b) What actions can cause transitions?
- c) Draw a small transition diagram.





# Environment for a Self-Driving Car

- ☐ **Fully observable:** The agent's sensors always show the whole **state**.

vs.

- ☐ **Partially observable:** The agent only perceives part of the **state** and needs to remember or infer the rest.

## Deterministic:

- ☐ a) **Percepts** are 100% reliable
- ☐ b) Changes in the environment are completely determined by the current **state** of the environment and the agent's **action**.

vs.

## Stochastic:

- ☐ a) **Percepts** are unreliable (noise distribution, sensor failure probability, etc.). This is called a stochastic sensor model.
- ☐ b) The **transition function** is stochastic leading to transition probabilities and a Markov process.

- ☐ **Known:** The agent knows the **transition function**.

vs.

- ☐ **Unknown:** The agent needs to **learn the transition function** by trying actions.



Check what applies and explain what it means for a self-driving car.





# What Type of Intelligent Agent is a Self-Driving Car?

☐ Is it learning?

☐ Utility-based agents

☐ Goal-based agents

☐ Model-based reflex agents

☐ Simple reflex agents

Does it collect utility over time? How would the utility for each state be defined?

Does it have a goal state?

Does it store state information. How would they be defined (atomic/factored)?

Does it use simple rules based on the current percepts?



Check what applies

# Why is this so hard?

- Self-driving cars operate in a very complicated partially observable, stochastic, and dynamic environment.
- Can only use bounded rationality because of limits with sensors and computational power.
- Require a set of different agents that cooperate.

